

École polytechnique de Louvain

Improving the performance and the scalability of INGINIOUS

Author: **Guillaume EVERARTS DE VELP**
Supervisor: **Ramin SADRE**
Readers: **Olivier BONAVENTURE, Anthony GÉGO**
Academic year 2019–2020
Master [120] in Computer Science

Acknowledgements

Guillaume Everarts de Velp, 2020

Abstract

Contents

1	Introduction	4
1.1	INGInious	4
1.1.1	Architecture	4
1.1.2	Key features	5
1.1.3	Bottlenecks	5
1.2	Scaling	6
1.3	Intentions	7
2	State of the art	8
2.1	Virtualization vs Containerization	8
2.2	Containerization solutions	8
2.2.1	Docker	8
2.2.2	Podman	8
2.2.3	runC	8
2.2.4	SAND	8
2.2.5	SOCK	8
2.3	Virtualization solutions	8
2.3.1	Firecracker	8
2.3.2	Kata Containers	8
2.3.3	LightVM	8
3	Measurements	9
4	Results	10
5	Conclusion	11

Chapter 1

Introduction

1.1 INGINious

INGInious is a web platform developped by the UCL. It is a tool for automatic correction of programs written by students. It currently relies on Docker containers to provide a good isolation between the machine hosting the site and the execution of the student's codes. So that a problem in a program submitted by a student couldn't have any impact on the platform. Docker also allow to manage the resources granted to each code execution. For now INGINious can meet the demand and provide honest performances and responsiveness. But looking at the growing usage of the platform, we might soon come to a point where we reach the limits of the current implementation.

1.1.1 Architecture

INGInious counts four main components:

1. The front-end: the website with which each student interacts when submitting a task.
2. The back-end: a queue of all the tasks that need to be graded.
3. The docker agent: responsible for the container assignment to the pending tasks when resources are available.
4. The docker containers: one for the student code and one for the teacher tests evaluating the student's code behaviour.

The journey of a task submitted on INGINious is represented on Figure 1.1.

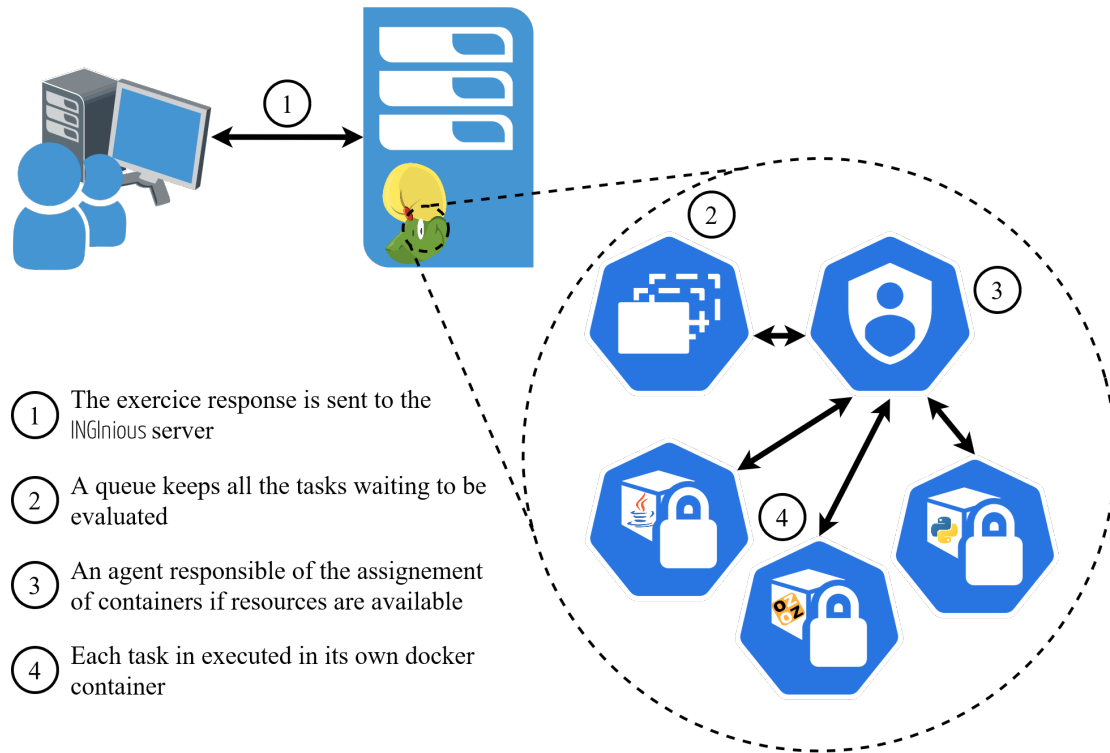


Figure 1.1: INGINIOUS global architecture

1.1.2 Key features

The key features of INGINIOUS, that allow it to meet the requirement of a code grading platform are the following:

- Isolation between the student's code and the platform.
- Resource limitation (CPU, RAM, Network) for the student's code execution.
- Modularity and versatility regarding the tasks that INGINIOUS can correct. Multiple programming language are supported, and new ones can easily be added.

1.1.3 Bottlenecks

When a task is submitted, we can count five delays before the answer can be delivered to the student:

- The sending time: the time it takes for the task to be sent to the back-end.

- The waiting time: the time the task will spend in the queue, waiting for an available container.
- The booting time: the time it takes to the container to boot and be ready to evaluate the task.
- The grading time: the time it takes for the code to run and for the teacher's container to grade it.
- The response time: the time to send the reponse to the student.

For the first and the last one, supposing that the machine hosting the website is not overwhelmed, the delay depends entirely on the network, this is a bit out of our hands here. The waiting time is directly related to the current load on the platform, this is a more a symptom of the server overwhelming than its cause, it could be directly solved by using a scaling strategies (see section 1.2). And then we come to the booting time and the grading time, which directly depends on the containerization technology used and on the hardware performances.

1.2 Scaling

Currently, the resources provided to INGIInious vary depending on the load that the platform is expected to be facing. Typically, when the grading of an exam is done by INGIInious, the platform is scaled up, and during the holidays it is scaled down. When it comes to scaling, two strategies can be used; vertical scaling and horizontal scaling.

Vertical scaling consists in adding more resources on a single machine, to improve its performances when needed. For example, when the number of tasks arriving to the server grow, we could increase the number of virtual Cores allocated to the Virtual Machine hosting the platform in order to be able to threat more of them concurrently. If the size of the waiting queue is increasing, we might want to make more RAM available.

Horizontal scaling consists in sharing the workload across multiple machines, so that each machine can handle a small part of it. This is a solution widely used nowadays as it allows to scale up virtually indefinetely, which is not the case with vertical scally where we depend on the maximum capacity of the hardware. This requires to rethink the architecture of the platform globally.

1.3 Intentions

The master thesis aims at improving INGINious, regarding its performances and its scalability. To do so, I will search and compare different containerization technologies that could be used instead of Docker. The goal is to find an alternative that decrease the booting time (and the grading time) without losing any of the key features of the platform. If such an alternative is found and proven to be worth the change, INGINious could then be refreshed with it.

Chapter 2

State of the art

2.1 Virtualization vs Containerization

2.2 Containerization solutions

2.2.1 Docker

2.2.2 Podman

2.2.3 runC

2.2.4 SAND

2.2.5 SOCK

2.3 Virtualization solutions

2.3.1 Firecracker

2.3.2 Kata Containers

2.3.3 LightVM

Chapter 3

Measurements

Chapter 4

Results

Chapter 5

Conclusion

Bibliography

UNIVERSITÉ CATHOLIQUE DE LOUVAIN

École polytechnique de Louvain

Rue Archimède, 1 bte L6.11.01, 1348 Louvain-la-Neuve, Belgique | www.uclouvain.be/epl