

École polytechnique de Louvain

Improving the performance and the scalability of INGINIOUS

Author: **Guillaume EVERARTS DE VELP**

Supervisor: **Ramin SADRE**

Readers: **Olivier BONAVENTURE, Anthony GÉGO**

Academic year 2019–2020

Master [120] in Computer Science and Engineering

Acknowledgements

Guillaume Everarts de Velp, 2020

Abstract

Contents

1	Introduction	5
1.1	INGInious	5
1.1.1	Architecture	5
1.1.2	Key features	6
1.1.3	Bottlenecks	6
1.2	Scaling	7
1.3	Intentions	8
2	Background knowledge	9
2.1	Virtualisation, Containerisation, Runtime isolation	9
2.1.1	Virtualisation	9
2.1.2	Containerisation	10
2.1.3	Runtime isolation	10
2.1.4	INGInious use case	11
2.2	Containers	11
2.2.1	Core concepts	11
2.2.2	Storage drivers	12
2.2.3	Container runtime	14
2.2.4	Container manager	15
2.3	Ecosystem overview	15
2.4	Containerisation solutions	16
3	Related work	19
3.1	Container and Virtual Machine performances	19
3.2	SAND	20
3.3	SOCK	20
3.4	LightVM	21
3.5	Firecracker	22
3.6	Summary	22
3.7	My master thesis	23

4	Benchmark tool	24
4.1	Setup	24
4.1.1	Testing environment	24
4.1.2	Configuration of the environment	25
4.2	Tests	25
4.2.1	Candidates	25
4.2.2	Experiments	27
5	Measurements	30
5.1	Results	30
5.2	Interpretation	30
5.3	Recommendations	30
5.4	Summary	30
6	Conclusion	31
A	Container life cycle	32

Chapter 1

Introduction

In this chapter I will simply introduce the subject of this master thesis, showing some basic concepts related to it and some key aspects.

1.1 INGINious

INGInious is a web platform developped by the UCL. It is a tool for automatic correction of programs written by students. It currently relies on Docker containers to provide a good isolation between the machine hosting the site and the execution of the student's codes. So that a problem in a program submitted by a student couldn't have any impact on the platform. Docker also allow to manage the resources granted to each code execution. For now INGINious can meet the demand and provide honest performances and responsiveness. But looking at the growing usage of the platform, we might soon come to a point where we reach the limits of the current implementation.

1.1.1 Architecture

INGInious counts four main components:

1. The front-end: the website with which each student interacts when submitting a task.
2. The back-end: a queue of all the tasks that need to be graded.
3. The docker agent: responsible for the container assignment to the pending tasks when resources are available.
4. The docker containers: one for the student code and one for the teacher tests evaluating the student's code behaviour.

The journey of a task submitted on INGINious is represented on Figure 1.1.

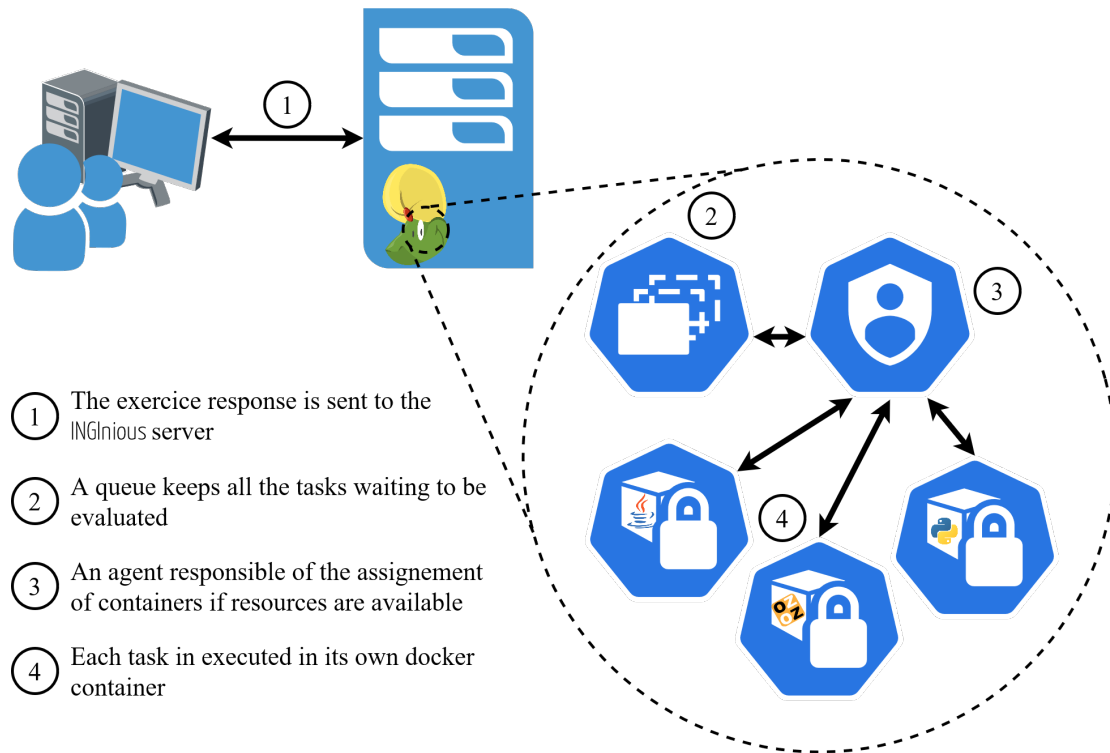


Figure 1.1: INGINious global architecture

1.1.2 Key features

The key features of INGINious, that allow it to meet the requirement of a code grading platform are the following:

- Isolation between the student's code and the platform.
- Resource limitation (CPU, RAM, Network) for the student's code execution.
- Modularity and versatility regarding the tasks that INGINious can correct. Multiple programming language are supported, and new ones can easily be added.

1.1.3 Bottlenecks

When a task is submitted, we can count five delays before the answer can be delivered to the student:

- The sending time: the time it takes for the task to be sent to the back-end.
- The waiting time: the time the task will spend in the queue, waiting for an available container.
- The booting time: the time it takes to the container to boot and be ready to evaluate the task.
- The grading time: the time it takes for the code to run and for the teacher's container to grade it.
- The response time: the time to send the response back to the student.

For the first and the last one, supposing that the machine hosting the website is not overwhelmed, the delay depends entirely on the network, this is a bit out of our hands here. The waiting time is directly related to the current load on the platform, this is a more a symptom of the server overwhelming than its cause, it could be directly solved by using a scaling strategies (see section 1.2). And then we come to the booting time and the grading time, which directly depends on the containerization technology used and on the hardware performances, this is where we are going to try to improve things in this thesis.

1.2 Scaling

Currently, the resources provided to INGINious vary depending on the load that the platform is expected to be facing. Typically, when the grading of an exam is done by INGINious, the platform is scaled up, and during the holidays it is scaled down. When it comes to scaling, two strategies can be used; vertical scaling and horizontal scaling.

Vertical scaling consists in adding more resources on a single machine, to improve its performances when needed. For example, when the number of tasks arriving to the server grow, we could increase the number of virtual Cores allocated to the Virtual Machine hosting the platform in order to be able to threat more of them concurrently. If the size of the waiting queue is increasing, we might want to make more RAM available.

Horizontal scaling consists in sharing the workload across multiple machines, so that each machine can handle a small part of it. This is a solution widely used nowadays as it allows to scale up virtually indefinitely, which is not the case with vertical scaling where we depend on the maximum capacity of the hardware. This requires to rethink the architecture of the platform globally.

1.3 Intentions

The master thesis aims at improving INGINious, regarding its performances and its scalability. To do so, I will search and compare different containerization technologies that could be used instead of Docker. The goal is to find an alternative that decreases the booting time (and the grading time) without losing any of the key features of the platform. If such an alternative is found and proven to be worth the change, INGINious could then be refreshed with it.

Chapter 2

Background knowledge

In this chapter I will put some basis, and explain some concepts related to my work. Reading this chapter should allow the reader to understand what I did in this work, regardless of its background.

2.1 Virtualisation, Containerisation, Runtime isolation

We always want the same thing, be able to execute some code, with no interaction with someone else's one. As of course we can not use one bare-metal machine for each application we want to run, some mechanism have to be used to provide isolation between coexisting processes on a same physical machine.

We can distinguish three level of isolation: Virtualisation, Containerisation and Runtime isolation (Figure 2.1).

2.1.1 Virtualisation

The isolation layer is located either between the hardware and the guest OS (hypervisor of type 1) or between a virtualisation of the hardware, by an hypervisor (type 2) running in another OS, and the guest OS. In both cases, the guest OS, will run as if it was on a bare-metal machine, executing its own kernel, managing its drivers, its file system, etc. This is the stronger isolation we can get. Cloud providers can rent you some Virtual Machines, this is what we call IaaS (Infrastructure as a Service). You can use many different hypervisor solutions, usually, Cloud providers have their very own solution.



Figure 2.1: Different isolation level for an application.

2.1.2 Containerisation

Here we go one level higher, the isolation layer is between the OS and the guest application. This isolation mechanism is called a container, and is basically a set of processes, running in common namespaces, with a root filesystem different from the host. Containers share the same kernel as the host though, which means that a process running in a container can be seen from the host. Containers belong to the world of PaaS (Platform as a Service), a system where Cloud providers offer you to run your application (that you provide under a containerized form) without having to worry about all the infrastructure that needs to be deployed along with it, and sometimes even with some great scalability mechanism support.

2.1.3 Runtime isolation

This is the highest level of (weakest) isolation you can get, the isolation is provided by the runtime on which the application is running (ex: The JVM for a Java program, a Python environment, ...). Multiple guest applications share the same OS, file system, without any mechanism to hide it from them. Cloud providers propose such service under the name of FaaS (Function as a Service). For this you can provide a small program, that you want to be executed on demand. Cloud providers usually allocate a container to this program, to provide a safer isolation,

but will execute multiple instances of this program in the same containers, to avoid the overhead of creating a new containers each time. The kinds of programs you usually run in these are computationnaly intensive parts of your application, that needs to have low response time and great scalability (for example resizing of user uploaded images). This model is often reference as the Lambda model (from Amazon Lambda service) or as serverless computing.

2.1.4 INGINious use case

The case of INGINious is a special one. The type of the workload caused by the tasks evaluation (fast response time, varying demand, potentially big computation, logic independent of the core application) would suggest to use a serverless strategy, but the isolation requirement are those of PaaS (as the various inputs of the functions are code of student to safely execute, isolated from one another). And for now the all application relies on IaaS, running in multiple VMs where are hosted the core application and multiple docker agents.

2.2 Containers

The isolation being the most important requirement, we will then look into containerisation solutions. Presenting the main concepts behind containerisation, the global container ecosystem, the current solution that INGINious uses and the alternatives we have.

2.2.1 Core concepts

Containers rely on three components: namespaces, control groups and chroot.

Namespaces are a feature of the Linux kernel since 2002, they are a key element that made containerisation possible. A namespace can be associated to a context, which is a partition of all the resources of a system that a set of processes has access to, while other processes can't. Those sets of resources can be of seven kinds:

- **mnt** (Mount): This controls the mount points. Processes can only have access to the mount point of their namespace.
- **pid** (Process ID): This allows each process in each namespace to get a process id assigned indepedently of other namespaces porcesses.
- **net** (Network): This provides a virtualized network stack.

- **ipc** (Interprocess Communication): This allows processes of a same namespace to communicate with one another, for example by sharing some memory.
- **uts** (Hostname): This allows to have different hostnames on a same machine, each hostname being considered as unique by the processes of its namespace.
- **user** (User ID): This allows to change the user id in a namespace.
- **cgroup** (Control group): This allows to change the root cgroup directory, this virtualize how process's cgroups are viewed.

When a Linux machine starts, it initiate one namespace of each type in which all the processes run. The processes can then choose to create and switch of namespaces.

Control groups are a feature of the Linux kernel that allows to limit and control the resources allocated to some processes. You can for example control the cpu usage, the memory consumption, the io... Recently (since kernel 4.5) a new version of cgroups (cgroups v2) appeared, which comes to tackle the flaws of the original implementation, while keeping all of its functionalities. Though, the adoption of the new version is a process, and takes times, and still now many applications use the original version of cgroup (like runc).

Chroot allows to change the root of the root filesystem for some processes. For example, we could create a directory `/tmp/myroot/` which contains all of the usual directories present in the original root folder (`/`) and set this as the new apparent root for the chosen processes. This is not a complete sandbox, and not a real isolation on its own, files from outside of the chroot could still be accessed.

2.2.2 Storage drivers

When it comes to handle a container file system, different solutions can be used. The goal of each is to provide the most efficient writable root directory (`/`) for each container, but keeping each container unaffected by the modification done in the other containers.

In order to do so, three main strategies can be used:

- **Deep copy**: for each container, the whole image is copied during the creation of the container. This is simple, but gets terribly slow as the file system size increases.

- **File based copy on write:** for each container, will be copied only the files that are edited during the container life cycle. This is more complex, but get more efficient as the container's size grows.
- **Block based copy on write:** for each container, will be copied only the blocks (in the filesystem) that are edited during the container life cycle. This is even more complex, but get more efficient as some small part of big files are edited.

Containers are a specific kind of workload in the sense that many information, data, is redundant in different containers. For example, for a simple application, we could use several containers with different responsibilities and tools embedded in it, but all based on the same Alpine image. This brought a new space problem, as we don't want to avoid duplicating too much data. In order to face this, **union filesystems** are used, along with layered container images. This basically means that different container images but with the same basis, will actually share this same basis, avoiding the need to duplicate it. It differs from file system deduplication mechanism by the level at which the action is taken. For deduplication, when a block is written on disk, we check if similar data isn't already located somewhere. Whereas for union file system, the common base image will be read-only, and when an attempt to modify a file is made, a bind mound will be made with a copy of the file on top of the original one, so that the common base stays preserved.

The storage of a container has to be backed by a file system, which will eventually already provide some interesting mechanisms for the container. Those four main filesystem used are:

- **BTRFS**, presented in 2013, [19] is a general purpose file system, based on copy-on-write and with efficient snapshots capabilities and strong data integrity.
- **zFS** is quite similar to BTRFS (but was there before), but handles some mechanisms differently, like the management of blocks (blocks vs. extends¹), snapshots (birth-times vs reference-counting²), ... It also support deduplication, which can be usefull for system were a lot of data is redundant and disk space is valuable.

¹zFS can have blocks to up to 128KB, while BTRFS will only use the extend strategy (point to the next block).

²When doing snapshots, zFS will store on the block the time at wich a new use of the block has been added, while Btrfs handles it with a reference-counting mechanism.

- **xFS** is a network file system (as zFS and BTRFS), this means that it can manage several drives in different locations, connected over the network in one storage unit. It manages to deliver better performances and availability than other similar solutions at the time it was created (1993). [20] It has no focus on any copy-on-write mechanism as the first two though.
- **ext4** was meant to be the replacement to ext3, as the "Linux Filesystem" when it was presented in 2007. [16] It aims at providing a good balance between scalability, reliability, performance and stability. It doesn't provide any of the fancy feature of the previous file system.

The Table 2.1 present a short summary of the different storage driver available for the different solutions.

Tool	Storage driver	C-o-W³	FS⁴
Docker/Podman	overlay2	file based	ext4, xFS
Docker/Podman	aufs	file based	ext4, xFS
Docker/Podman	devicemapper	block based	direct-lvm ⁵
Docker/Podman	btrfs	block based	BTRFS
Docker/Podman	zfs	block based	zFS
Docker/Podman	vfs	no	any
LXD	Btrfs	block based	BTRFS
LXD	ZFS	block based	zFS
LXD	LVM	yes	ext4
LXD	Directory	no	any

Table 2.1: Summary of storage driver solutions.

2.2.3 Container runtime

A container runtime is a tool that creates containers, executes process in it, and deletes dead containers. It will have the responsibility to create the namespaces, change the root directory, and attach processes to a control group. The most common container runtime nowadays is **runc**, created by the Open Container project.

³Copy-on-write

⁴Backing file system

⁵Note that this is a logical volume manager, not a file system, which uses in our case the xFS file system

2.2.4 Container manager

A container manager will use the container runtime, to provide a "user-friendly" interface to manage containers. It has the responsibility to set up the network interfaces, to provide the image of the containers and any Copy-on-write mechanism that could go along with it. It creates and manages the control groups in with containers will be set. Some of them offer the possibility to create custom images, to create pods, or swarm, which are entities of multiple interconnected containers.

2.3 Ecosystem overview

A small overview of the current container ecosystem can be found on Figure 2.2. Note that solutions like Kubernetes⁶ which are more oriented towards hosting and continuous deployment of container based applications than to single container provisioning are not presented here⁷.

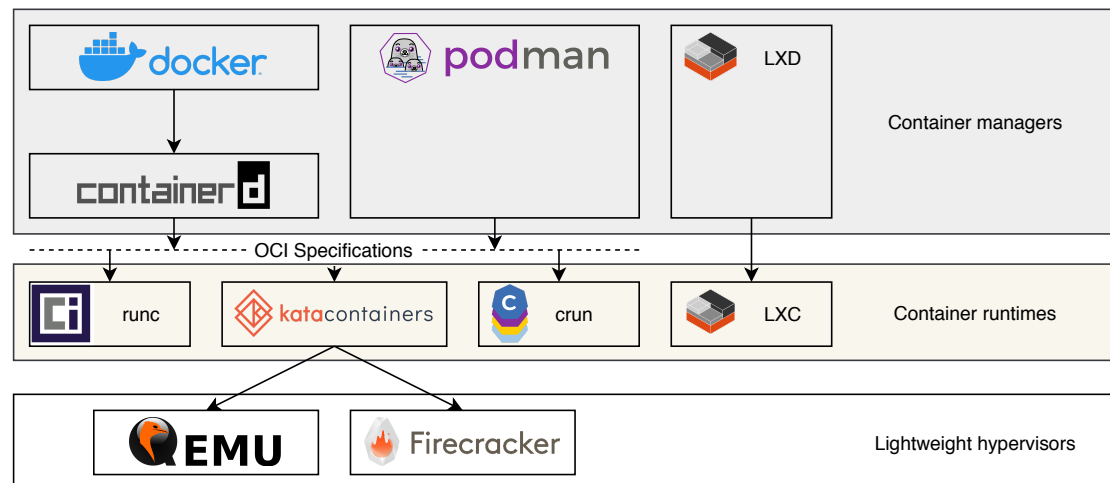


Figure 2.2: Small overview of the current container ecosystem

OCI (Open Container Initiative) is "an open governance structure for the express purpose of creating open industry standards around container formats and runtime." [3] They currently have two specifications: the format of the image that

⁶Kubernetes is "an open-source system for automating deployment, scaling, and management of containerized applications" [5]

⁷A more detailed overview for those kind of applications is provided by Containerd at the following address: <https://containerd.io/img/architecture.png>

has to be given to an OCI compatible runtime, and the commands to interact with such runtime.

2.4 Containerisation solutions

Docker [17] is today probably the most known solution for containerization. Since its apparition in 2014, its interest for the PaaS sector hasn't ceased to grow. It consists in a daemon, running on the host, that allows to easily manage different containers. It relies on Containerd, "An industry-standard container runtime with an emphasis on simplicity, robustness and portability" [1], which is a more complete container runtime than what I presented before, with embedded network management, storage driver management, and other mechanism. Containerd is not meant to be used standalone though, which is why we still use Docker. By default, Docker uses **runc** as container runtime by default, but is compatible with any runtime that fulfill the OCI [3] requirements.

The main advantages of Docker are its simplicity of use, its production-grade quality, the huge fleet of ready-to-run containers publicly available and a lot of useful tools (like **docker-swarm** that come along with it). This is the solution currently used by INGINious.

"**Podman** is a daemonless container engine for developing, managing, and running OCI [3] Containers on your Linux System." [6] Podman presents itself as a viable true alternative to Docker. Its main difference is the fact that it runs daemonless, it runs containers as detached child processes and containers can be rootless⁸ by default (while it is only a recent feature coming up to Docker). It can also manage pods, which are groups of containers deployed on the same machine. Its default runtime is **runc** as well. Podman is an Open-Source project, and is still growing a lot, the latest version at this day (2020-02-28) is v1.8.0, released 21 days ago, and already 287 commits have been done since then!

"**LXC** is a userspace interface for the Linux kernel containment features." [7] This solution is developed and maintained by Canonical Ltd. Docker used to be based on LXC until it created its own execution environment. LXC came out recently with a new solution; LXD, which offer about the same things as Docker does. A bunch of ready to run images publicly available, and a nice command line interface to interact with containers. The only missing feature of LXC compare to Docker, regarding our use case, is the possibility to launch a container with

⁸A rootless container is a container in which the processes running in the container are not owned by the root user of the host machine, and neither is the runtime managing the container.

a command, and stop it when it is finished. LXC actually start a complete init process for each container, in which you can then come and execute your command.

Linux-VServer [2] is a bit of a special one. No namespaces are used to provide isolation here, they rely on contexts to separate multiple virtualised systems. They have modified the Linux kernel to provide some utilities like process isolation, network isolation and CPU isolation. The main advantage of this approach is scalability as it can run a large amount of containers, the drawback is the portability. It requires a modified version of the kernel and it also misses some features like checkpoint and resume that we can have with more classical containerization and real virtualisation. This is an project much older than Docker and its contemporaries and differ in goal, it is not appropriate for the INGINious workload. This is an experience more similar to classical VM ones, but with an isolation provided by containerisation.

OpenVZ [4] is another container-based virtualisation solution. But this time they use namespaces to provide it. They also have their own way of dealing with system resources management, more complete than cgroup. Even though this more close to what runc and LXC so in term of isolation, this is still not the kind of experience appropriate for INGINious use case.

Kata Containers is not another container manager. It is an OCI compatible runtime. With the specificity that it doesn't run mainstream containers as **runc**, but actually lightweight virtual machines, using KVM virtualisation and with Qemu (originally), Firecracker or Cloud Hypervisor (still in its early days) as hypervisor. Those two first currently available hypervisor are not really equivalent though, as the original one has more features (like devices assignment) and the second one has lower memory footprint, smaller attack surface, and a shortest boot time. They put forward four features of their solution:

- **Security:** Thanks to their virtualization solution, each runtime has its own kernel, with real network, i/o and memory isolation.
- **Compatibility:** They support industry standarts, as the OCI [3] and legacy virtualization technologies.
- **Performance:** Their performances are consistent with classical containerization solutions.
- **Simplicity:** They eliminate the need to have a virtual machine dedicated to host containers.

crun is a complete equivalent of **runc**, yet another OCI compatible runtime, but this one is implemented in C, which give it a small performance advantage over **runc**, implemented in GO. **crun** has also full support for cgroupv2, which is still lacking for **runc** at the moment (2020-04-22).

Chapter 3

Related work

In this chapter I will present some of the work that has already been done in the field of containerization, and show what we can take from it. I will then explain why my work adds up something to those previous research.

3.1 Container and Virtual Machine performances

As said previously, Virtual Machines have been kept away from the PaaS world, because of the greater overhead that Virtualization has on booting time and resources consumptions of the isolated system created. This paper [11] made a detailed point about it back in 2014, emphasizing on three points that needed to be improved for the next generation of containers: Security, OS Independence and Standardization. The last one meaning that containers runtime should establish some common requirements, come up with a common container format. We have seen this appear since then, under the OCI (Open Container Initiative) specifications. A tradeoff with the performances (mostly I/O) of the application running isolated is much more evident with virtualisation than with containerisation [12].

But recently things started to change, some new solutions seem to have the advantage of both sides, the isolation of a VM and the portability and lightness of a Container. This is why some VM-based solution are presented here. Even though VMs do not come from the same domain of application as Containers, and even though in the case of INGINious, Containers were 5 years ago the obvious choice, it might now be time to reconsider it.

3.2 SAND

SAND [9] aims at improving the performance of serverless computing with two techniques: application-level sandboxing, and a hierarchical message bus:

Application Sandboxing: The key idea here is to differentiate multiple executions of different functions with multiple execution of a same function. In the first case, as usual, a new container is launched on the new function demand and the function is executed inside of it. In the second case, instead of launching a new container with exactly the same configuration as an already running one, we only fork a process inside the running container to deal with the new demand, which is much faster than creating a new container. The problem in our case, is that we lose the isolation between student code execution, which is not desirable.

Hierarchical Message Queuing: The goal to this is to facilitate communication between functions that interact with one another (i.e. the output of a function is the input of another one). To do so they use a two-level communication bus: global and local. Global means that the communication is made between different functions on different hosts, while local means different functions on the same host. As accessing the local bus is much faster than the global one, it can decrease latency for functions on the same host. For now in the case of INGINIOUS, student and teacher container always run on the same host and communicate through the means of file descriptors. So this is still not something for us.

3.3 SOCK

"Sock (roughly for serverless-optimized containers), [is] a special-purpose container system with two goals: (1) *low-latency invocation* for Python handlers that import libraries and (2) *efficient sandbox initialization* so that individual workers can achieve high ready-state throughput." [18] The final product created here isn't really something we could use for INGINIOUS, it is a serverless solution (based on the Lambda model), targeting specifically python applications. Though, in the process of creating this solution, they started from containers and deconstructed their performances, identifying their bottlenecks. And from this we can take those things:

- Bind mounting is twice as fast as AUFS. Even though AUFS (used by Docker) as a useful Copy-on-write capability, we don't need to write to most of our files in our case, so using a read-only bind mounting for those files could allow us to avoid copying all files before startup, while still being able to edit the one we want to.

- Network namespaces creation and cleanup are costly, due to a single lock shared across all network namespaces. We might gain some performances by not adding network interfaces to containers that don't need it. This is actually already done by the docker agent of INGINious.
- Reusing a cgroup is at least twice as fast as creating a new one each time. We could keep a pool of initialized cgroups, and only change the current container its controls.

3.4 LightVM

LightVM is a complete redesign of Xen's toolstack which tried to bring some container's characteristics to VMs. Such as fast instantiation (small startup time) and high instance density capability (high number of instances running in parallel on the same machine). [15] Xen is a Type-1¹ hypervisor presented in 2003 with really low virtualization overheads and high hosting capacity, allowing a machine to host up to a 100 guest OS. [10]

They achieve such container's like performances by:

- reducing the image size and the memory footprint of virtual machines. They do so by including in the VM only what is necessary to the application that is meant to be executed in it.
- introducing noxs (no XenStore), a new implementation of Xen, without XenStore (which was a real bottleneck for fast instantiation of multiple VMs).
- splitting the Xen's toolstack into what can be run before the VM creation and what has to be done during it. Allowing to pre-initialize VMs.
- replacing the Hotplug Script by xendevd, a binary daemon that can execute pre-defined setup more efficiently.

They present four use cases with this solution, one of them being more interesting for us: lightweight computation services, for which they rely on Minipython unikernel, to run computations written in Python, as a Faas could propose to do. This would still be hard to use for INGINious, for the same reason as for SOCK (§3.3).

¹A Type-1 hypervisor is an hypervisor that runs on a bare-metal machine, without additional host.

3.5 Firecracker

Firecracker is a new VMM (hypervisor) created specifically for serverless and containers applications. [8] This is a solution provided by AWS, that very recently got deployed for two of their web services: Lambda (Faas) and Fargate (Paas). We are basically getting here the good isolation of virtual machines and nearly as good performances and low overhead of containers. Firecracker is based on KVM and provides minimal virtual machines (MicroVMs). The configuration is done through a REST API. Device emulation is available for disks, networking and serial console. Network can be limited and so can disk throughput and request rate. If proven to be easily usable in INGINious case, this would provide a better alternative to classical containerisation regarding security.

3.6 Summary

In this section, on Table 3.1, are quickly reminded the several solutions we explored in this chapter, along with some interesting infos about them.

Name	Pub.	Update	Open-Source	Pot. Sol.	Com.	Isol.
SAND	2018	?	?	No	No	Cont.
SOCK	2018	?	?	No	No	Cont. or Runt.
LigthVM	2017	2017	Yes	No	No	Virt.
Firecracker	2019	2020	Yes	Yes	Yes	Virt.

Table 3.1: Summary table of the different solutions explored in this chapter.

Caption:

- *Name*: The name of the project.
- *Pub.*: The first publication year of the project.
- *Update*: The last update year of the project.
- *Open-source*: If the project is open-source.
- *Pot. Sol.*: If the project could be a solution in our case.
- *Com.*: If the project is a "commercial grade" solution.

- *Isol.*: The type of isolation used in the project.
- *Cont.*: Isolation by containerization.
- *Virt.*: Isolation by virtualization.
- *Runt.*: Isolation by the runtime of the application.

3.7 My master thesis

There is already been a huge amount of performance comparison done for the containerisation world, essentially comparing bare-metal machine to containers and VMs.

- *Performance evaluation of container-based virtualization for high performance computing environments* [21] (2013): They compared a native experience to LXC containers, OpenVZ, VServer and Xen. They showed that virtualisation (with Xen) was giving really poor performances compared to the others.
- *Performance overhead comparison between hypervisor and container based virtualization* [14] (2017): In this paper they showed that even though container-based virtualization (with Docker) is more lightweight than real virtualisation, its performances are not always better. And they plan to compare it into more details later.
- *A performance comparison of container-based technologies for the cloud* [13] (2017): They compared here two different containerization solutions (Docker and LXC) against the native experience with focus on CPU usage, memory consumption, network and I/O performances.

In short they all show that containers perform in general better than VMs and worse than a Physical machine. What they don't compare is different containerisations solution with one another (or at least not all of them) and they don't focus on booting time. This is the direction I will take in this report. Trying to respond to the question: "What would be the most appropriate containerisation (or virtualisation) solution for INGINious use case?", or in other word, how to get the shortest instantiation time, along with the best isolation possible and as low overall performance overhead as we can get.

Chapter 4

Benchmark tool

In this chapter I will present the setup used to perform my measurements, the different tests I made and the candidate solutions that are going to be compared with one another.

4.1 Setup

4.1.1 Testing environment

Originally, all of my tests were meant to be executed in a Virtual Machine, on my computer. In a VM because this is currently the nature of the environment where INGINIOUS launches its container. And on my computer, because it is powerful enough for the modest load of tests that I am running for each solution. Unfortunately, the first setup I had didn't support nested virtualisation (which is required for Kata-based solutions), the second one was very unstable with it and I finally moved to the cloud, renting Azure VMs.

The early results that I could get were quite surprising though, showing really poor performances for Kata containers with Firecracker, and it turned out that Firecracker is not really meant to be running inside VMs (it can work, but not at its best). I made then a last move, to an Intel NUC provided by the INGI department. All the tests will then run on a bare-metal machine, no more nested virtualization problem at the horizon.

The testing environment used for the final results presented in this work is an Intel NUC with the following configuration:

processor	Intel(R) Core(TM) i7-7567U CPU @ 3.50GHz
memory	16GB
main partition	2TB HDD
additionnal storage	500GB SSD
operation system	Ubuntu 18.04.4 LTS
linux kernel version	4.15.0-99-generic

Note about storage: The main partition, on which the operating system is running, is a hard drive. This will provide much lower IO performances than an SSD. To be more in pair with the kind of performances we can get in the cloud, all the relevant data accessed and written during my tests will be stored on the additionnal disk (SSD).

4.1.2 Configuration of the environment

As a lot of different solutions are going to be tested, and all of them with their own requirements (sometimes conflicting with the ones of other solutions), some work has to be done to make it easier for someone to replicate them. The configuration of the environment is going to be done with Ansible¹ playbooks, with one playbook for each solution to test.

All the different configurations allong with the tests they are required for and their corresponding playbook are presented in section 4.2.1.

4.2 Tests

As we already said previously, time is an important aspect in the experience INGINious provides to its users. The time for a task to be evaluated and the time before this task is evaluated. In order to improve those two, we will focus in those tests on the booting time of containers and their IO handling (still related to the time overhead that some solution could have compared to another).

While we are at it, we will also verify if some solutions handle network much poorly than others. Each test done is presented in more details in section 4.2.2.

4.2.1 Candidates

A candidate solution is a combinaison of different elements, variabilities, choices to make. We will consider here those ones:

¹<https://docs.ansible.com/>

- **Container manager:** *Docker*, *Podman*, *LXD*, those are the different user friendly solution that can be used to manage container, and that we will compare here.
- **Container runtime:** *runc*, *crun*, *LXC*, *Kata containers* (Qemu or Firecracker), the (less user-friendly) container runtimes, taking care of all the isolation that a container require.
- **Control group:** *cgroup* and *cgroupv2*, there is no real choice to make here, *cgroupv2* is the successor of *cgroup*, and should be used, when supported by the other element of the configuration.
- **Storage driver:** *overlay*, *aufs*, *btrfs*, *zfs*, *vfs*, *lvm*, those are the different strategies that can be used to manage the file system of the container.
- **Base container image:** *Alpine* because it is the default choice when it comes to conceive containerized applications today and *Centos* because it is the current choice made by INGINious.
- **Rootless container:** Run the container without root privileges, neither of the container runtime, the container manager and the processes running in the container are executed as root. This feature has to be supported by the container manager and the container runtime, and provide better security.

Though, we can not compose a candidate solution with one element of each category, randomly picked. Some element of the solution might only be usable with some of the possibilities for a variability. The different constraints that we have are listed here:

- Docker doesn't support *cgroupv2* yet (actually the support has to be added by *containerd*).
- LXD doesn't support *cgroupv2* yet.
- LXD only supports LXC as runtime, and LXC is only supported by LXD.
- LXD only supports *btrfs*, *zfs*, *vfs* (under the name of *dir*) and *lvm* as storage driver.
- Kata Container with Firecracker only supports *devicemapper* (*lvm*) as storage driver.
- *Zfs*, *lvm*, *aufs* can not be used with rootless containers.
- *Cgroup* (v1) can not be used with rootless containers.

The list of different candidate solution we can compare is presented in Table 4.1. Each container will be launched with a CPU limitation of one single core, and memory limitation of 1GB. This is more than enough for each test that will be done. A network interface will be only added to the container if it needs it for the specific test.

#	Cont. Mngr	Base Image	Storage driver	Ctrl grp	Runtime
1	Docker	Alpine	overlay2	cgroup	runc
2	Docker	Alpine	aufs	cgroup	runc
3	Docker	Alpine	devicemapper	cgroup	runc
4	Docker	Alpine	btrfs	cgroup	runc
5	Docker	Alpine	zfs	cgroup	runc
6	Docker	Alpine	vfs	cgroup	runc
7	Docker	Centos	overlay2	cgroup	runc
8	Docker	Alpine	overlay2	cgroup	kata-qemu
9	Docker	Alpine	devicemapper	cgroup	kata-qemu
10	Docker	Alpine	devicemapper	cgroup	kata-fc
11	Docker	Alpine	overlay2	cgroup	crun
12	Podman	Alpine	overlay2	-	runc
13	Podman	Alpine	overlay2	cgroupv2	crun
14	-	Alpine	-	cgroup	runc
15	-	Alpine	-	cgroupv2	crun
16	LXD	Alpine	btrfs	cgroup	LXC
17	LXD	Alpine	zfs	cgroup	LXC
18	LXD	Alpine	directory	cgroup	LXC
19	LXD	Alpine	lvm	cgroup	LXC

Table 4.1:

4.2.2 Experiments

For all the experiments presented here, time measurements are taken at four important steps of the execution of the container:

t_0 Before its creation, this is time zero.

t_c After its creation. Container's state is *created*.

t_s After its startup. Container's state is *running*.

t_e After the end of the execution of the command. Container's state is *running*.

This gives us three interesting measurements:

1. Creation time: $t_{create} = t_c - t_0$
2. Starting time: $t_{start} = t_s - t_c$
3. Execution time: $t_{exec} = t_e - t_s$

In order to do this, each container will be running as first process a shell², which will never receive any input, but is here to ensure that the container will stay up as long as we need it. When we don't need the container anymore, we simply kill all the processes in it.

Booting time

This tests aims at determining the candidate solution that can provide a ready to run container the faster. In order to measure this, a simple container composed of only the base image is created, in which the simple command `/bin/echo Hello World` is executed. The best candidate would be the one that has the shortest creation and starting time ($t_{create} + t_{start}$).

Basic I/O performances

This tests aims at determining mostly the influence of the usage of each storage driver on the performances of an application that relies highly on it. Four different tests are done here: big file read and write, and great amount of file read and write. The best candidate solution would be the one to with the shortest creation, starting and execution time ($t_{create} + t_{start} + t_{exec}$).

For the first one, five sqlite databases have been created, of five different size (151.6 kB, 536.6 kB, 2.6 MB, 11.9 MB and 111.6 MB, about one order of magnitude appart from each other). We use an sqlite database as everything will be stored in one file, and no daemon needs to be runnning, which reduces the performance cost of each solicitation to purely I/O (and eventually CPU, if the requests were to be complex, which is not our case here). The read test will do a select operation on each table of the database without any filter, displaying all of the table content. The write test will duplicate each table in the database. The database chosen as basis to generate all the presented one is `tpcc`³, which a benchmark database, perfect for our case. The biggest database used is an export

²except for LXD, which natively has an init process running in each container

³<https://relational.fit.cvut.cz/dataset/TPCC>

from the tpcc database, the other ones are created by removing content of some tables from this same database.

For the second one, five loads of small file have been created (10, 100, 1000, 10 000 and 100 000 files). The content of the file is alphanumeric, and completely random (no file should be a duplication of another one), each file is 2048 character (and bytes) long. This represent the best the kind of file a user could access, to read or write content from/to it. The read test will load the content of each file in memory, then discard it and move to the next one. The write test will extract an uncompressed archive (`tar`) containing all those files.

Network setup time

This tests aims also at determining the candidate solution that can setup a ready du run container the fastest, but with the added complexity of adding a network interface, and mapped ports, so that the host machine can send requests to an http server running inside of the container. I choosed a `lighttpd` server, for its lightness, and its wide use in container application. We take one more time measurement here, the time at which we received the first response at the http requests we send to our server (t_r). The best candidate would be the one that has the shortest delay before this response ($t_{create} + t_{start} + t_{exec} + t_{response}$, with $t_{response} = t_r - t_e$). The execution time corresponds to the execution of the command that launches the server as detached process.

Ping response time

This really simple test aims at determining if some solution provide a much weaker network solution to the container than the others. We simply perform a ping request to `1.1.1.1`, and compare the response time we get for each solution. The best solution would be the one with the shortest ping response time.

Chapter 5

Measurements

5.1 Results

5.2 Interpretation

5.3 Recommendations

5.4 Summary

Chapter 6

Conclusion

Appendix A

Container life cycle

During its life cycle, a container will go through five stages: creation, launch, execution, stopping and cleaning. Those five steps are explained below and presented on figure A.1, along with the corresponding command for each presented tool. Other optional stages (like pause/unpause) are not presented here.

1. **Create** This is the set-up phase, depending on the tool, the file system of the new container might be copied, or any other things that need to be done before launching the container. This phase can only be done one time by container.
2. **Launch** This is the proper instantiation of the container, the namespaces are created, the new file system is adopted. Depending on the container, a complete init process might be executed. This state can be repeated as many times as we want for a container, as long as it is in a stopped state.
3. **Execute** This corresponds to the execution of a process in the container. This can be done as many times as we want for a container, as long as the container is running.
4. **Stop** This is when the container needs to be stopped, all running processes are killed, and the namespaces are exited. No modification should be done on the filesystem, the storage of the container remains.
5. **Clean** This is the un-create phase, where we delete everything that was done during the creation phase.




Figure A.1: Container life cycle and how to interact with it.

Bibliography

- [1] An industry-standard container runtime with an emphasis on simplicity, robustness and portability. <https://containerd.io/>. Accessed: 2020-02-26.
- [2] Linux vserver. <http://linux-vserver.org/Paper>. Accessed: 2020-04-25.
- [3] Open container initiative. <https://www.opencontainers.org/>. Accessed: 2020-03-04.
- [4] Openvz. <https://openvz.org>. Accessed: 2020-04-25.
- [5] Production-grade container orchestration. <https://kubernetes.io/>. Accessed: 2020-02-26.
- [6] What is podman? <https://podman.io/whatis.html>. Accessed: 2020-02-28.
- [7] What's lxc? <https://linuxcontainers.org/fr/lxc/introduction/#whats-lxc>. Accessed: 2020-03-01.
- [8] Alexandru Agache, Marc Brooker, Andreea Florescu, Alexandra Iordache, Anthony Liguori, Rolf Neugebauer, Phil Piwonka, and Diana-Maria Popa. Firecracker: Lightweight virtualization for serverless applications.
- [9] Istemi Ekin Akkus, Ruichuan Chen, Ivica Rimac, Manuel Stein, Klaus Satzke, Andre Beck, Paarijaat Aditya, and Volker Hilt. Sand: Towards high-performance serverless computing. In *2018 USENIX Annual Technical Conference (USENIX ATC 18)*, pages 923–935, 2018.
- [10] Paul Barham, Boris Dragovic, Keir Fraser, Steven Hand, Tim Harris, Alex Ho, Rolf Neugebauer, Ian Pratt, and Andrew Warfield. Xen and the art of virtualization. *ACM SIGOPS operating systems review*, 37(5):164–177, 2003.
- [11] Rajdeep Dua, A Reddy Raja, and Dharmesh Kakadia. Virtualization vs containerization to support paas. In *2014 IEEE International Conference on Cloud Engineering*, pages 610–614. IEEE, 2014.

- [12] Wes Felter, Alexandre Ferreira, Ram Rajamony, and Juan Rubio. An updated performance comparison of virtual machines and linux containers. In *2015 IEEE international symposium on performance analysis of systems and software (ISPASS)*, pages 171–172. IEEE, 2015.
- [13] Zhanibek Kozhirkbayev and Richard O Sinnott. A performance comparison of container-based technologies for the cloud. *Future Generation Computer Systems*, 68:175–182, 2017.
- [14] Zheng Li, Maria Kihl, Qinghua Lu, and Jens A Andersson. Performance overhead comparison between hypervisor and container based virtualization. In *2017 IEEE 31st International Conference on Advanced Information Networking and Applications (AINA)*, pages 955–962. IEEE, 2017.
- [15] Filipe Manco, Costin Lupu, Florian Schmidt, Jose Mendes, Simon Kuenzer, Sumit Sati, Kenichi Yasukata, Costin Raiciu, and Felipe Huici. My vm is lighter (and safer) than your container. In *Proceedings of the 26th Symposium on Operating Systems Principles*, pages 218–233, 2017.
- [16] Avantika Mathur, Mingming Cao, Suparna Bhattacharya, Andreas Dilger, Alex Tomas, and Laurent Vivier. The new ext4 filesystem: current status and future plans. In *Proceedings of the Linux symposium*, volume 2, pages 21–33. Citeseer, 2007.
- [17] Dirk Merkel. Docker: lightweight linux containers for consistent development and deployment. *Linux journal*, 2014(239):2, 2014.
- [18] Edward Oakes, Leon Yang, Dennis Zhou, Kevin Houck, Tyler Harter, Andrea Arpaci-Dusseau, and Remzi Arpaci-Dusseau. Sock: Rapid task provisioning with serverless-optimized containers. In *2018 USENIX Annual Technical Conference (USENIX ATC 18)*, pages 57–70, 2018.
- [19] Ohad Rodeh, Josef Bacik, and Chris Mason. Btrfs: The linux b-tree filesystem. *ACM Transactions on Storage (TOS)*, 9(3):1–32, 2013.
- [20] Randolph Y Wang and Thomas E Anderson. xfs: A wide area mass storage file system. In *Proceedings of IEEE 4th Workshop on Workstation Operating Systems. WWOS-III*, pages 71–78. IEEE, 1993.
- [21] Miguel G Xavier, Marcelo V Neves, Fabio D Rossi, Tiago C Ferreto, Timoteo Lange, and Cesar AF De Rose. Performance evaluation of container-based virtualization for high performance computing environments. In *2013 21st Euromicro International Conference on Parallel, Distributed, and Network-Based Processing*, pages 233–240. IEEE, 2013.

UNIVERSITÉ CATHOLIQUE DE LOUVAIN
École polytechnique de Louvain

Rue Archimède, 1 bte L6.11.01, 1348 Louvain-la-Neuve, Belgique | www.uclouvain.be/epl