

École polytechnique de Louvain

Improving the performance and the scalability of INGINIOUS

Author: **Guillaume EVERARTS DE VELP**
Supervisor: **Ramin SADRE**
Readers: **Olivier BONAVENTURE, Anthony GÉGO**
Academic year 2019–2020
Master [120] in Computer Science

Acknowledgements

Guillaume Everarts de Velp, 2020

Abstract

Contents

1	Introduction	4
1.1	INGInious	4
1.1.1	Architecture	4
1.1.2	Key features	5
1.1.3	Bottlenecks	5
1.2	Scaling	6
1.3	Intentions	7
2	State of the art	8
2.1	Virtualization vs Containerization	8
2.2	Serverless	9
2.3	Namespaces	10
2.4	Containerization solutions	10
2.4.1	Docker	10
2.4.2	Podman	11
2.4.3	LXC	11
2.4.4	SAND	11
2.4.5	SOCK	12
2.5	Virtualization solutions	12
2.5.1	Firecracker	13
2.5.2	Kata Containers	13
2.5.3	LightVM	13
2.6	Summary	13
3	Measurements	14
4	Results	15
5	Conclusion	16

Chapter 1

Introduction

In this chapter I will simply introduce the subject of this master thesis, showing some basic concepts related to it and some key aspects.

1.1 INGINious

INGInious is a web platform developped by the UCL. It is a tool for automatic correction of programs written by students. It currently relies on Docker containers to provide a good isolation between the machine hosting the site and the execution of the student's codes. So that a problem in a program submitted by a student couldn't have any impact on the platform. Docker also allow to manage the resources granted to each code execution. For now INGINious can meet the demand and provide honest performances and responsiveness. But looking at the growing usage of the platform, we might soon come to a point where we reach the limits of the current implementation.

1.1.1 Architecture

INGInious counts four main components:

1. The front-end: the website with which each student interacts when submitting a task.
2. The back-end: a queue of all the tasks that need to be graded.
3. The docker agent: responsible for the container assignment to the pending tasks when resources are available.
4. The docker containers: one for the student code and one for the teacher tests evaluating the student's code behaviour.

The journey of a task submitted on INGINious is represented on Figure 1.1.

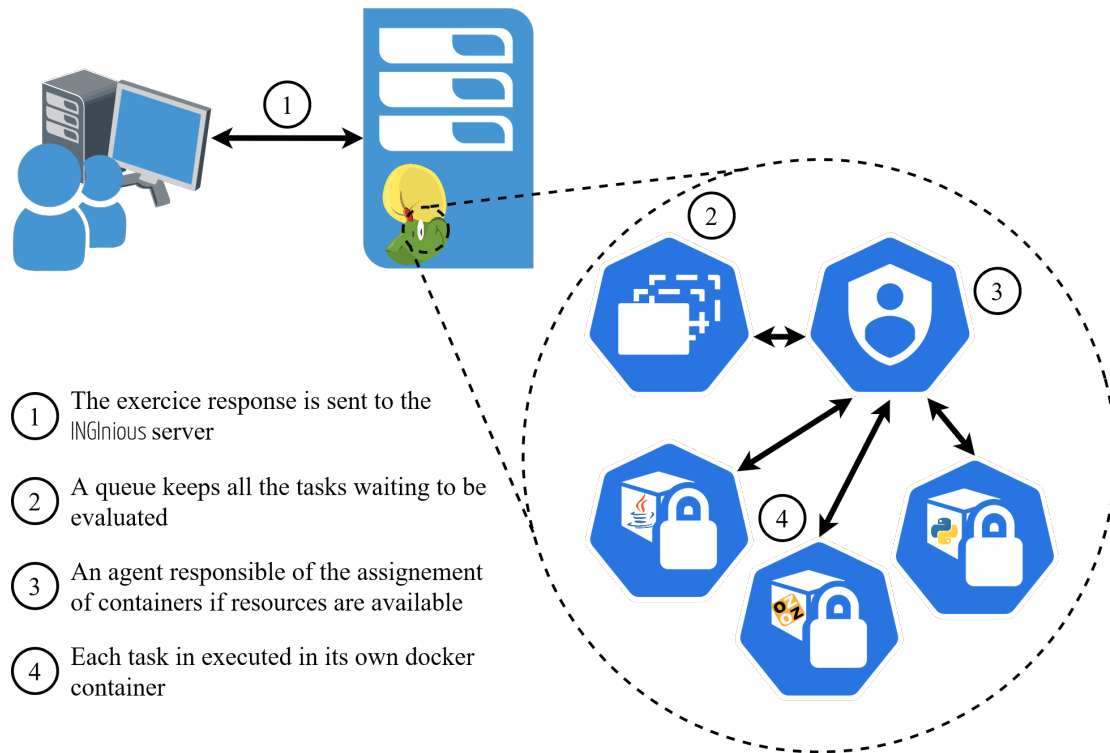


Figure 1.1: INGINious global architecture

1.1.2 Key features

The key features of INGINious, that allow it to meet the requirement of a code grading platform are the following:

- Isolation between the student's code and the platform.
- Resource limitation (CPU, RAM, Network) for the student's code execution.
- Modularity and versatility regarding the tasks that INGINious can correct. Multiple programming language are supported, and new ones can easily be added.

1.1.3 Bottlenecks

When a task is submitted, we can count five delays before the answer can be delivered to the student:

- The sending time: the time it takes for the task to be sent to the back-end.
- The waiting time: the time the task will spend in the queue, waiting for an available container.
- The booting time: the time it takes to the container to boot and be ready to evaluate the task.
- The grading time: the time it takes for the code to run and for the teacher's container to grade it.
- The response time: the time to send the response back to the student.

For the first and the last one, supposing that the machine hosting the website is not overwhelmed, the delay depends entirely on the network, this is a bit out of our hands here. The waiting time is directly related to the current load on the platform, this is a more a symptom of the server overwhelming than its cause, it could be directly solved by using a scaling strategies (see section 1.2). And then we come to the booting time and the grading time, which directly depends on the containerization technology used and on the hardware performances, this is where we are going to try to improve things in this thesis.

1.2 Scaling

Currently, the resources provided to INGINious vary depending on the load that the platform is expected to be facing. Typically, when the grading of an exam is done by INGINious, the platform is scaled up, and during the holidays it is scaled down. When it comes to scaling, two strategies can be used; vertical scaling and horizontal scaling.

Vertical scaling consists in adding more resources on a single machine, to improve its performances when needed. For example, when the number of tasks arriving to the server grow, we could increase the number of virtual Cores allocated to the Virtual Machine hosting the platform in order to be able to threat more of them concurrently. If the size of the waiting queue is increasing, we might want to make more RAM available.

Horizontal scaling consists in sharing the workload across multiple machines, so that each machine can handle a small part of it. This is a solution widely used nowadays as it allows to scale up virtually indefinitely, which is not the case with vertical scaling where we depend on the maximum capacity of the hardware. This requires to rethink the architecture of the platform globally.

1.3 Intentions

The master thesis aims at improving INGINious, regarding its performances and its scalability. To do so, I will search and compare different containerization technologies that could be used instead of Docker. The goal is to find an alternative that decreases the booting time (and the grading time) without losing any of the key features of the platform. If such an alternative is found and proven to be worth the change, INGINious could then be refreshed with it.

Chapter 2

State of the art

In this chapter I will present how things are today in the field of containerization. I will start with a quick reminder of the difference between virtualization and containerization, then go through some existing solutions for both of those concepts.

2.1 Virtualization vs Containerization

Both virtualization and containerization aims at providing an abstraction layer that allows the application or the OS (Operating System) located above the abstraction to behave like if it was the only one present on this layer. The key difference between those two is where this abstraction layer is located.

To keep it simple, for the virtualization, the abstraction is between the OS and the hardware, the OS thinks it is running on a baremetal machine but is actually hosted on another machine. For the containerization, the abstraction is between the application and the OS, the application thinks it is the only one running in the system but it is actually only isolated in its own namespace. Multiple containers on a same host share then the same kernel. As a good illustration is better than a thousand words, Figure 2.1 represents this difference.

Because of what they are, those two solutions usually didn't apply to the same situations. A VM (Virtual Machine) is a much more complicated and complete structure, that takes more time to boot but provides a better isolation. And the Container is much lighter, quicker to launch but less isolated. Traditionally, in the Cloud Computing world, VMs rely in IaaS (Infrastructure-as-a-Service) while Containers rely in PaaS (Platform-as-a-Service). This paper [8] made a detailed point about it back in 2014.

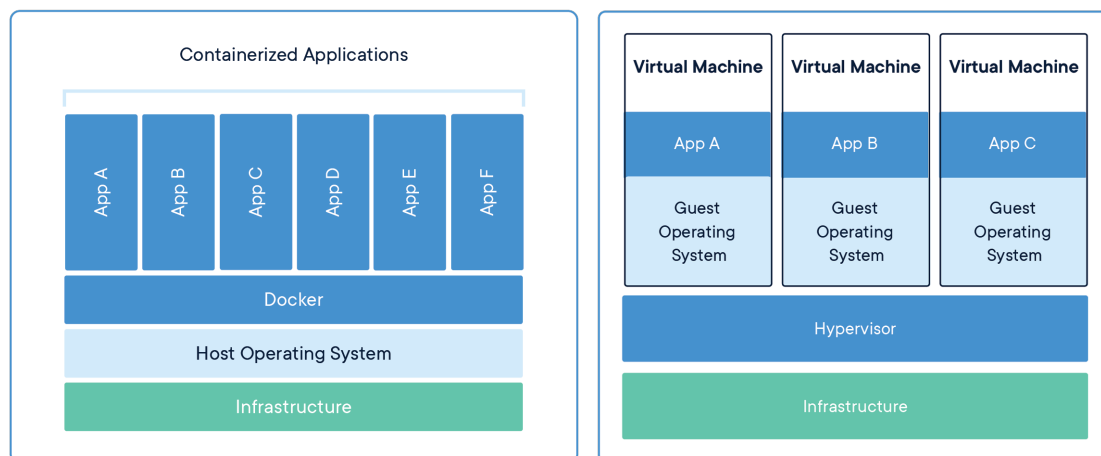


Figure 2.1: This image comes from Docker’s website [4].

But recently things started to change, some new solutions seem to have the advantage of both sides, the isolation of a VM and the portability and lightness of a Container. This is why some VM-based solution are presented here. Even though VMs do not come from the same domain of application as Containers, and even though in the case of INGINious, Containers were 5 years ago the obvious choice, it might now be time to reconsider it.

2.2 Serverless

Serverless can be seen as an improvement of Paas, with better scalability and efficiency. This is often called Faas (Function as a service), two open-source solutions for this are Apache OpenWhisk and OpenLambda. It can really still be associated with what INGINious does when testing a submission; it is basically, running a function, on user demand, with no direct link with the structure of the application. The main difference between what INGINious does and what serverless offers is that usually, in serverless, you run a bunch of functions over and over again, only with different inputs. With INGINious, the different inputs to those functions are the code of the student to execute, which makes the function execution unique for each submission. The function as a service style is often referenced as the Lambda model (from Amazon’s Lambda). Compared to a container, the isolation here will be even higher, as the same runtime is usually shared between multiple execution of the function. [9] This makes serverless solutions not really appropriate for INGINious use case, as we really *need* to provide an isolation between student’s tasks execution.

2.3 Namespaces

Namespaces are a feature of the Linux kernel since 2002, they are a key element that made containerization possible. A namespace can be associated to a context, which is a partition of all the resources of a system that a set of processes has access to, while other processes can't. Those sets of resources can be of seven kinds:

- **mnt** (Mount): This controls the mount points. Processes can only have access to the mount point of their namespace.
- **pid** (Process ID): This allows each process in each namespace to get a process id assigned independently of other namespaces processes.
- **net** (Network): This provides a virtualized network stack.
- **ipc** (Interprocess Communication): This allows processes of a same namespace to communicate with one another, for example by sharing some memory.
- **uts** (Hostname): This allows to have different hostnames on a same machine, each hostname being considered as unique by the processes of its namespace.
- **user** (User ID): This allows to pretend to change the user in a namespace, so that it could behave like he had more privileges for example.
- **cgroup** (Control group): This provides a virtualization of the resources available to the processes of the namespace.

When a Linux machine starts, it initiates one namespace of each type in which all the processes run. The processes can then choose to create and switch of namespaces.

2.4 Containerization solutions

Here below are presented some containerization solutions.

2.4.1 Docker

Docker [11] is today probably the most known solution for containerization. Since its apparition in 2014, its interest for the PaaS sector hasn't ceased to grow. It is now tightly binded with Kubernetes, "an open-source system for automating deployment, scaling, and management of containerized applications" [2]. It consists in a daemon, running on the host, that allows to easily manage different containers.

It relies on Containerd, "An industry-standard container runtime with an emphasis on simplicity, robustness and portability" [1], which itself uses runc ("runc is a CLI tool for spawning and running containers according to the OCI specification." [3]).

The main advantages of Docker are its simplicity of use, its production-grade quality, the huge fleet of ready-to-run containers publicly available and a lot of useful tools (like `docker-compose` that come along with it). This is the solution currently used by INGINious.

2.4.2 Podman

"Podman is a daemonless container engine for developing, managing, and running OCI Containers on your Linux System." [5] Podman present itself as a viable true alternative to Docker. Its main difference is the fact that it runs daemonless, it runs containers as child processes, it isn't therefore based on containerd as docker is. It can also manage pods, which are groups of containers deployed on the same machine. Podman is an Open-Source project, and is still growing a lot, the latest version at this day (2020-02-28) is v1.8.0, released 21 days ago, and already 287 commits have been done since then!

2.4.3 LXC

"LXC is a userspace interface for the Linux kernel containment features." [6] This solution is developed and maintained by Canonical Ltd. Docker used to be based on LXC until it created its own execution environment. The main difference between LXC and Docker is that LXC launches a complete Linux init for each new container, while Docker will only launch a service, the application, in the container.

2.4.4 SAND

SAND [7] aims at improving the performance of serverless computing with two techniques: application-level sandboxing, and a hierarchical message bus:

- **Application Sandboxing:** The key idea here is to manage differentiate multiple executions of different functions with multiple execution of a same function. In the first case, as usual, a new container is launched on the new function demand and the function is executed inside of it. In the second case, instead of launching a new container with exactly the same configuration as an already running one, we only fork a process inside the running container to deal with the new demand, which is much faster than creating a new container. The problem in our case, is that we lose the isolation between student code execution, which is not desirable.

- **Hierarchical Message Queuing:** The goal to this is to facilitate communication between functions that interact with one another (i.e. the output of a function is the input of another one). To do so they use a two-level communication bus: global and local. Global means that the communication is made between different functions on different hosts, while local means different functions on the same host. As accessing the local bus is much faster than the global one, it can decrease latency for functions on the same host. In our case, as INGINious is still running on a single host, this shouldn't improve our performances that much. But in a near future where INGINious get a new horizontally scaled architecture, it might be interesting to come back on it.

2.4.5 SOCK

"Sock (roughly for serverless-optimized containers), [is] a special-purpose container system with two goals: (1) *low-latency invocation* for Python handlers that import librairies and (2) *efficient sandbox initialization* so that individual workers can achive high ready-state throughput." [12] The final product created here isn't really something we could use for INGINious, it is a serverless solution (based on the Lambda model), targeting specifically python applications. Though, in the process of creating this solution, they started from containers and deconstructed their perfromances, indentifying their bottlenecks. And from this we can take those things:

- Bond mounting is twice as fast as AUFS. Even though AUFS (used by Docker) as a usefull Copy-on-write capability, we don't need to write to most of our files in our case, so using a read-only bind mounting for those files could allow us to avoid copying all file before startup, while still beeing able to edit the one with want to.
- Network namespaces creation and cleanup are costly, due to a single lock shared across all network namespaces. We might gain some performances buy not adding network interfaces to containers that don't need it.
- Reusing a cgroup is at least twice as fast as creating a new one each time. We could keep a pool of initialized cgroups, and only change the current container it controls.

2.5 Virtualization solutions

Here bellow are presented some lightweight virtualization solutions.

2.5.1 Firecracker

2.5.2 Kata Containers

2.5.3 LightVM

[10]

2.6 Summary

Chapter 3

Measurements

Chapter 4

Results

Chapter 5

Conclusion

Bibliography

- [1] An industry-standard container runtime with an emphasis on simplicity, robustness and portability. <https://containerd.io/>. Accessed: 2020-02-26.
- [2] Production-grade container orchestration. <https://kubernetes.io/>. Accessed: 2020-02-26.
- [3] runc's github repository. <https://github.com/opencontainers/runc#introduction>. Accessed: 2020-02-28.
- [4] What is a container? <https://www.docker.com/resources/what-container>. Accessed: 2020-02-24.
- [5] What is podman? <https://podman.io/whatis.html>. Accessed: 2020-02-28.
- [6] What's lxc? <https://linuxcontainers.org/fr/lxc/introduction/#whats-lxc>. Accessed: 2020-03-01.
- [7] Istemi Ekin Akkus, Ruichuan Chen, Ivica Rimac, Manuel Stein, Klaus Satzke, Andre Beck, Paarijaat Aditya, and Volker Hilt. Sand: Towards high-performance serverless computing. In *2018 USENIX Annual Technical Conference (USENIX ATC 18)*, pages 923–935, 2018.
- [8] Rajdeep Dua, A Reddy Raja, and Dharmesh Kakadia. Virtualization vs containerization to support paas. In *2014 IEEE International Conference on Cloud Engineering*, pages 610–614. IEEE, 2014.
- [9] Scott Hendrickson, Stephen Sturdevant, Tyler Harter, Venkateshwaran Venkataramani, Andrea C Arpaci-Dusseau, and Remzi H Arpaci-Dusseau. Serverless computation with openlambda. In *8th USENIX Workshop on Hot Topics in Cloud Computing (HotCloud 16)*, 2016.
- [10] Filipe Manco, Costin Lupu, Florian Schmidt, Jose Mendes, Simon Kuenzer, Sumit Sati, Kenichi Yasukata, Costin Raiciu, and Felipe Huici. My vm is lighter (and safer) than your container. In *Proceedings of the 26th Symposium on Operating Systems Principles*, pages 218–233, 2017.

- [11] Dirk Merkel. Docker: lightweight linux containers for consistent development and deployment. *Linux journal*, 2014(239):2, 2014.
- [12] Edward Oakes, Leon Yang, Dennis Zhou, Kevin Houck, Tyler Harter, Andrea Arpaci-Dusseau, and Remzi Arpaci-Dusseau. Sock: Rapid task provisioning with serverless-optimized containers. In *2018 USENIX Annual Technical Conference (USENIX ATC 18)*, pages 57–70, 2018.

UNIVERSITÉ CATHOLIQUE DE LOUVAIN
École polytechnique de Louvain

Rue Archimède, 1 bte L6.11.01, 1348 Louvain-la-Neuve, Belgique | www.uclouvain.be/epl